

qg_2172_HW4_Part_123

November 25, 2019

1 ELEN 6885 Reinforcement Learning Coding Assignment (Part 1, 2, 3)

1.1 Taxi Problem Overview

Please put your code into the block marked by: ##### YOUR CODE STARTS HERE YOUR CODE ENDS HERE ##### You should not edit anything outside of the block.

2 Playing with the environment

Run the cell below to get a feel for the environment by moving your agent(the taxi) by taking one of the actions at each step.

```
[7]: from gym.wrappers import Monitor
import gym
import random
import numpy as np
from google.colab import files

uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving utils.py to utils.py

Saving evaluation_utils.py to evaluation_utils.py

```
[0]: """
    You can test your game now.
    Input range from 0 to 5:
        0 : South (Down)
        1 : North (Up)
        2 : East (Right)
        3 : West (Left)
        4: Pick up
        5: Drop off
        6: exit_game
```

```

"""
GAME = "Taxi-v3"
env = gym.make(GAME)
env = Monitor(env, "taxi_simple", force=True)
s = env.reset()
steps = 100
for step in range(steps):
    env.render()
    action = int(input("Please type in the next action:"))
    if action==6:
        break
    s, r, done, info = env.step(action)
    print('state:',s)
    print('reward:',r)
    print('Is state terminal?:',done)
    print('info:',info)

# close environment and monitor
env.close()

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

```

↳
↳ -----
KeyboardInterrupt                                Traceback (most recent call↳
↳ last)

  /usr/local/lib/python3.6/dist-packages/ipykernel/kernelbase.py in ↳
↳ _input_request(self, prompt, ident, parent, password)
    729             try:
--> 730                 ident, reply = self.session.recv(self.stdin_socket,↳
↳ 0)
    731             except Exception:

  /usr/local/lib/python3.6/dist-packages/jupyter_client/session.py in ↳
↳ recv(self, socket, mode, content, copy)
    802             try:

```

```

--> 803             msg_list = socket.recv_multipart(mode, copy=copy)
      804             except zmq.ZMQError as e:

/usr/local/lib/python3.6/dist-packages/zmq/sugar/socket.py in
↳recv_multipart(self, flags, copy, track)
      465         """
--> 466         parts = [self.recv(flags, copy=copy, track=track)]
      467         # have first part already, only loop while more to receive

zmq/backend/cython/socket.pyx in zmq.backend.cython.socket.Socket.recv()

zmq/backend/cython/socket.pyx in zmq.backend.cython.socket.Socket.recv()

zmq/backend/cython/socket.pyx in zmq.backend.cython.socket._recv_copy()

/usr/local/lib/python3.6/dist-packages/zmq/backend/cython/checkrc.pxd in
↳zmq.backend.cython.checkrc._check_rc()

```

KeyboardInterrupt:

During handling of the above exception, another exception occurred:

```

KeyboardInterrupt                                Traceback (most recent call
↳last)

<ipython-input-2-5951a1d297f9> in <module>()
      17 for step in range(steps):
      18     env.render()
--> 19     action = int(input("Please type in the next action:"))
      20     if action==6:
      21         break

/usr/local/lib/python3.6/dist-packages/ipykernel/kernelbase.py in
↳raw_input(self, prompt)
      703         self._parent_ident,
      704         self._parent_header,
--> 705         password=False,
      706     )

```

```

/usr/local/lib/python3.6/dist-packages/ipykernel/kernelbase.py in
↪ _input_request(self, prompt, ident, parent, password)
733         except KeyboardInterrupt:
734             # re-raise KeyboardInterrupt, to truncate traceback
--> 735             raise KeyboardInterrupt
736         else:
737             break

```

```
KeyboardInterrupt:
```

2.1 1.1 Incremental implementation of average

We've finished the incremental implementation of average for you. Please call the function to estimate with $1/\text{step}$ step size and fixed step size to compare the difference between these two on a simulated Bandit problem.

```

[0]: def estimate(OldEstimate, StepSize, Target):
    '''An incremental implementation of average.
    OldEstimate : float
    StepSize : float
    Target : float
    '''
    NewEstimate = OldEstimate + StepSize * (Target - OldEstimate)
    return NewEstimate

```

```

[0]: import random
import numpy as np
random.seed(6885)
numTimeStep = 10000
q_h = np.zeros(numTimeStep + 1) # Q Value estimate with 1/step step size
q_f = np.zeros(numTimeStep + 1) # Q value estimate with fixed step size
FixedStepSize = 0.5 #A large number to exaggerate the difference
for step in range(1, numTimeStep + 1):
    if step < numTimeStep / 2:
        r = random.gauss(mu = 1, sigma = 0.1)
    else:
        r = random.gauss(mu = 3, sigma = 0.1)

    #TIPS: Call function estimate defined in ./RLalgs/utils.py
    #####
    # YOUR CODE STARTS HERE
    q_h[step]=estimate(q_h[step-1],1/step,r)

```

```

q_f[step]=estimate(q_f[step-1],FixedStepSize,r)

# YOUR CODE ENDS HERE
#####

q_h = q_h[1:]
q_f = q_f[1:]

```

Plot the two Q value estimates. (Please include a title, labels on both axes, and legends)

```

[0]: import matplotlib.pyplot as plt
      %matplotlib inline
      #####
      # YOUR CODE STARTS HERE
      plt.title('Q Value Estimates of fixed step size and dynamic step size')

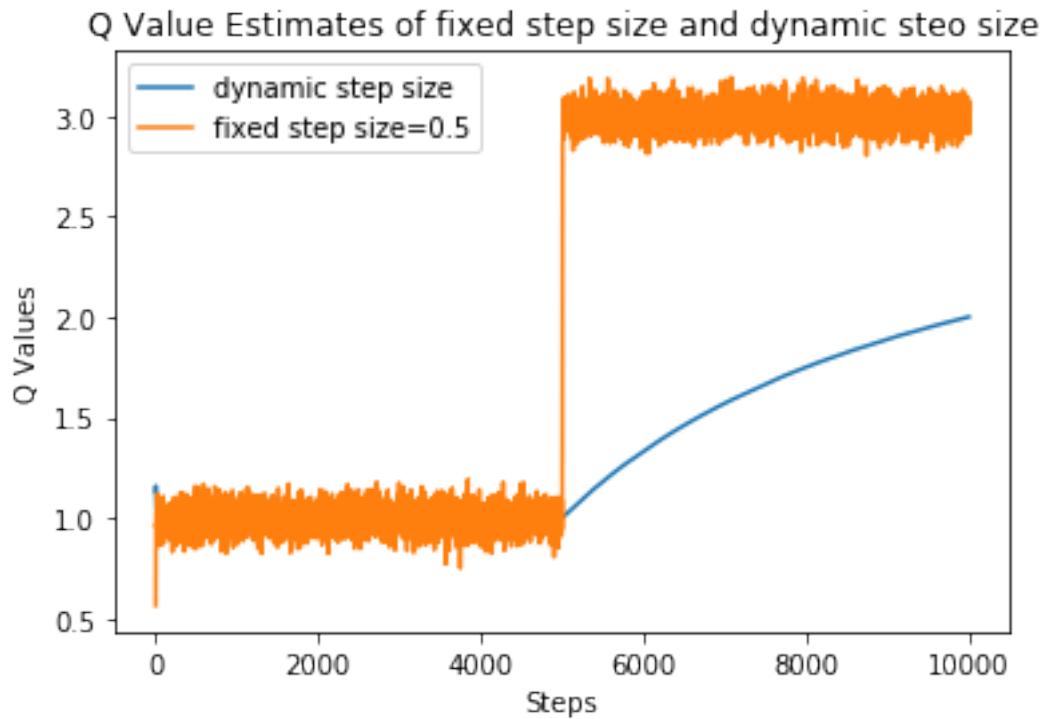
      plt.plot(q_h, label="dynamic step size")
      plt.plot(q_f, label="fixed step size=0.5")
      plt.legend()

      plt.xlabel('Steps')
      plt.ylabel('Q Values')

      plt.show()

      # YOUR CODE ENDS HERE
      #####

```



2.2 1.2 ϵ -Greedy for Exploration

In Reinforcement Learning, we are always faced with the dilemma of exploration and exploitation. ϵ -Greedy is a trade-off between them. You are supposed to implement Greedy and ϵ -Greedy. We combine these two policies in one function by treating Greedy as ϵ -Greedy where $\epsilon = 0$. Edit the function `epsilon_greedy` the following block.

```
[0]: def epsilon_greedy(value, e, seed = None):
    '''
    Implement Epsilon-Greedy policy.

    Inputs:
    value: numpy ndarray
    A vector of values of actions to choose from
    e: float
    Epsilon
    seed: None or int
    Assign an integer value to remove the randomness

    Outputs:
    action: int
    Index of the chosen action
    '''
    assert len(value.shape) == 1
```

```

assert 0 <= e <= 1

if seed != None:
    np.random.seed(seed)

#####
# YOUR CODE STARTS HERE
probability=np.full(value.shape,e/value.shape[0])
max_value_index=np.argmax(value)
probability[max_value_index]+=1-e
action=np.random.choice(value.shape[0],1,p=probability)

# YOUR CODE ENDS HERE
#####

return action

```

```

[0]: np.random.seed(6885) #Set the seed for reproducibility
q = np.random.normal(0, 1, size = 5)
#####
# YOUR CODE STARTS HERE
greedy_action=epsilon_greedy(q,0, None)
e_greedy_action=epsilon_greedy(q,0.1, None)

# YOUR CODE ENDS HERE
#####
print('Values:')
print(q)
print('Greedy Choice =', greedy_action)
print('Epsilon-Greedy Choice =', e_greedy_action)

```

Values:

```
[ 0.61264537  0.27923079 -0.84600857  0.05469574 -1.09990968]
```

Greedy Choice = [0]

Epsilon-Greedy Choice = [0]

You should get the following results: Values: [0.61264537 0.27923079 -0.84600857 0.05469574 -1.09990968] Greedy Choice = 0 Epsilon-Greedy Choice = 0

2.3 1.3 Exploration VS. Exploitation

Try to reproduce Figure 2.2 (the upper one is enough) of the Sutton's book based on the experiment described in Chapter 2.3.

```

[0]: # # Do the experiment and record average reward acquired in each time step
# #####
# # YOUR CODE STARTS HERE

```

```

# Returns the action-value for each action at the current time step
def Q_cal(actions):
    array_size=len(actions)
    Q=[] # used to store Q value for every state

    # action[i][0] represents total reward of action i
    # action[i][1] represents number of actions taken of action i
    # for example action[3][10,5] means for action 5, it has been taken 5 times,
    →with a total reward of 10

    # update Q
    for i in range(array_size):
        if actions[i][1]==0:
            Q.append(0) # If this action hasn't been taken yet, return 0
        else:
            Q.append(actions[i][0]/float(actions[i][1]))# total reward/ number of
    →action
    return Q

def epsilon_greedy_iteration(values,e,iteration_num):
    array_size=len(values)
    actions=[]
    rewards=[]

    # initialize every action's Q value
    for i in range(array_size):
        actions.append([0.0,0])#( [0,0.0],[0,0.0],[0.0.0]...)

    for j in range(iteration_num):
        r=np.random.rand(1) # random reward
        # greedy policy: epsilon=e
        if r>e:
            action_to_take=np.argmax(Q_cal(actions)) # greedy
        else:
            action_to_take=np.random.randint(0,array_size) # equal probability for
    →every action

        reward=np.random.normal(values[action_to_take],scale=1.0,size=1)[0] #
    →arbitrary reward~(value,1)
        actions[action_to_take][0]+=reward # uodate Q value of this action
        actions[action_to_take][1]+=1

        # record each action's average reward for each step
        rewards.append(reward)
    return rewards # return a 1*n convex that store reward for each step (with
    →different actions)

```



```

def run(e,steps):
    reward_total=0
    rewards=[]
    for i in range(2000): # to obtain a stable value
        values=np.random.normal(0,1,size=10)
        reward=epsilon_greedy_iteration(values,e,steps)
        rewards.append(reward)
    means=np.mean(rewards,axis=0) # return a 1*n vector that stores mean reward
    →from step 1 to step n after trying 2000 times
    return means
# # YOUR CODE ENDS HERE
# #####


```

```

[0]: # Plot the average reward
#####
# YOUR CODE STARTS HERE
import matplotlib.pyplot as plt
%matplotlib inline
greedy=run(0.0,1000)
epsilon_001 = run(0.01, 1000)
epsilon_010 = run(0.1, 1000)
plt.plot(range(0, 1000), greedy,label='greedy')
plt.plot(range(0, 1000), epsilon_001,label='epsilon = 0.01')
plt.plot(range(0, 1000), epsilon_010,label='epsilon = 0.1')
plt.xlabel('Steps')
plt.ylabel('Average reward')
plt.title('Comparison')
plt.legend()
plt.show()

# YOUR CODE ENDS HERE
#####

```



3 Question 2

In this question, you will implement the value iteration and policy iteration algorithms to solve the Taxi game problem

3.1 2.1 Model-based RL: value iteration

For this part, you need to implement the helper functions `action_evaluation(env, gamma, v)`, and `extract_policy(env, v, gamma)` in `utils.py`. Understand `action_selection(q)` which we have implemented. Use these helper functions to implement the `value_iteration` algorithm below.

```
[0]: import numpy as np
import utils
def value_iteration(env, gamma, max_iteration, theta):
    """
    Implement value iteration algorithm. You should use extract_policy to for_
    ↪ extracting the policy.

    Parameters
    -----
    env: OpenAI env.
        env.P: dictionary
            the transition probabilities of the environment
```

```

        P[state][action] is tuples with (probability, nextstate,
→reward, terminal)
        env.nS: int
            number of states
        env.nA: int
            number of actions

gamma: float
    Discount factor.
max_iteration: int
    The maximum number of iterations to run before stopping.
theta: float
    Determines when value function has converged.
Returns:
-----
value function: np.ndarray
policy: np.ndarray
"""
V = np.zeros(env.nS, dtype=float)
#####
# YOUR CODE STARTS HERE
policy=np.zeros(env.nS, dtype=int)
for i in range(max_iteration):
    Vprime=np.zeros(env.nS, dtype=float)

    for state in range(env.nS):
        for action in range(env.nA):
            V_tmp=0
            for probability, nextstate, reward, terminal in env.
→P[state][action]: # directly find the maximum V value for each state
                V_tmp=V_tmp+probability*(reward+gamma*V[nextstate])
            if (Vprime[state]<V_tmp):
                Vprime[state]=V_tmp # update real-time maximum V value for
→each state
            if np.all(np.abs(V-Vprime)<theta): # judge when to converge and break
→if so
                break
    V=Vprime.copy() # update V

# Find optimal policy by given optimal V
new_policy=np.zeros(env.nS, dtype=int)
for state in range(env.nS):
    current_max_v=-999999
    tmp=np.zeros(env.nA, dtype=float) # save real-time V value for one state
→according to different actions
    for action in range(env.nA):
        for probability, nextstate, reward, terminal in env.P[state][action]:

```

```

        tmp[action]=tmp[action]+probability*(reward+gamma*V[nextstate]) #
→calculate current V for current action
        if(tmp[action]>current_max_v):# judge whether it is the maximum V
→value for this state until now
            current_max_v=tmp[action] # If so, update ir
            new_policy[state]=action # If so, update policy

    policy=new_policy.copy()

    # YOUR CODE ENDS HERE
    #####

    return V, policy

```

After implementing the above function, read and understand the functions implemented in `evaluation_utils.py`, which we will use to evaluate our value iteration policy

```

[0]: import evaluation_utils
import gym
GAME = "Taxi-v3"
env = gym.make(GAME)
V_vi, policy_vi = value_iteration(env, gamma=0.95, max_iteration=6000,
→theta=1e-5)
# visualize how the agent performs with the policy generated from value
→iteration
evaluation_utils.render_episode(env, policy_vi)

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

(South)

```

+-----+
|R: | : :G|
| : | : : |

```

```
| : | : : : | |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(West)

```
+-----+
```

```
|R: | : :G| |
| : | : : |
| | : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(West)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : : : |
| | | : | : |
|Y| : |B: |
```

```
+-----+
```

(South)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(South)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(Pickup)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : : : |
| | | : | : |
|Y| : |B: |
```

```
+-----+
```

(North)

```
+-----+
```

```
|R: | : :G|
| : | : : |
```

```
|█: : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(North)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : █: : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(East)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : █: : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(East)

```
+-----+
```

```
|R: | : :G|
| : | █: : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(North)

```
+-----+
```

```
|R: | █: :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(North)

```
+-----+
```

```
|R: | : █:G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(East)

```
+-----+
```

```
|R: | : :G|
| : | : : |
```

```

| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(East)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(Dropoff)
Episode reward: 6.000000

```

```
[0]: # evaluate the performance of value iteration over 100 episodes
evaluation_utils.avg_performance(env, policy_vi)
```

```
[0]: 8.282828282828282
```

3.2 2.2 Model-based RL: policy iteration

In this part, you are supposed to implement policy iteration to solve the Taxi game problem.

```
[0]: import utils
def policy_iteration(env, gamma, max_iteration, theta):
    """Implement Policy iteration algorithm.

    You should use the policy_evaluation and policy_improvement methods to
    implement this method.

    Parameters
    -----
    env: OpenAI env.
        env.P: dictionary
            the transition probabilities of the environment
            P[state][action] is tuples with (probability, nextstate,
→reward, terminal)
        env.nS: int
            number of states
        env.nA: int
            number of actions
    gamma: float
        Discount factor.
    max_iteration: int
        The maximum number of iterations to run before stopping.
    theta: float
        Determines when value function has converged.
```

```

Returns:
-----
value function: np.ndarray
policy: np.ndarray
"""

V = np.zeros(env.nS, dtype=int)
policy = np.zeros(env.nS, dtype=int)

for state in range(env.nS):
    policy[state]=1 # intial policy for every state
    #####
    # YOUR CODE STARTS HERE
    for i in range(max_iteration):
        Vprime=policy_evaluation(env,policy,gamma,theta) # policy evaluation
        policy_improved,flag=policy_improvement(env,Vprime,policy,gamma) #
→policy improvement
        if np.all(np.abs(V-Vprime)<theta): # judge when to converge and break
→if so
            break
        V=Vprime.copy() # update V value for each state
        policy=policy_improved.copy() # update policy for each state

    # YOUR CODE ENDS HERE
    #####

return V, policy

def policy_evaluation(env, policy, gamma, theta):
    """Evaluate the value function from a given policy.

    Parameters
    -----
    env: OpenAI env.
        env.P: dictionary
            the transition probabilities of the environment
            P[state][action] is tuples with (probability, nextstate,
→reward, terminal)
        env.nS: int
            number of states
        env.nA: int
            number of actions

    gamma: float
        Discount factor.

```



```

policy: np.array
    The policy to evaluate. Maps states to actions.
max_iteration: int
    The maximum number of iterations to run before stopping.
theta: float
    Determines when value function has converged.
Returns
-----
value function: np.ndarray
    The value function from the given policy.
"""
V = np.zeros(env.nS)

#####
# YOUR CODE STARTS HERE
Vprime=np.zeros(env.nS) # value of next state

while True:
    Vprime=np.zeros(env.nS,dtype=float)
    for state in range(env.nS):
        action=policy[state] # choose action according to policy
        for probability, nextstate, reward, terminal in env.
→P[state][action]:
            ↱
→Vprime[state]=Vprime[state]+probability*(reward+gamma*V[nextstate]) # update ↱
→V value for each state
        if np.all(np.abs(V-Vprime)<theta): # judge when to converge and break ↱
→if so
            break
    V=Vprime.copy() # update current V

# YOUR CODE ENDS HERE
#####

return V

def policy_improvement(env, value_from_policy, policy, gamma):
    """Given the value function from policy, improve the policy.

    Parameters
    -----
    env: OpenAI env
        env.P: dictionary
            the transition probabilities of the environment

```

```

        P[state][action] is tuples with (probability, nextstate,
→reward, terminal)
        env.nS: int
            number of states
        env.nA: int
            number of actions

value_from_policy: np.ndarray
    The value calculated from the policy
policy: np.array
    The previous policy.
gamma: float
    Discount factor.

Returns
-----
new_policy: np.ndarray
    An array of integers. Each integer is the optimal action to take
    in that state according to the environment dynamics and the
    given value function.
stable_policy: bool
    True if the optimal policy is found, otherwise false
"""
#####
# YOUR CODE STARTS HERE
new_policy=np.zeros(env.nS,dtype=int)

policy_stable=True

for state in range(env.nS):
    current_max_v=-999999
    tmp=np.zeros(env.nA,dtype=float) # save real-time V value for one state
→according to different actions
    for action in range(env.nA):
        for probability,nextstate,reward,terminal in env.P[state][action]:
            □
→tmp[action]=tmp[action]+probability*(reward+gamma*value_from_policy[nextstate])
→# calculate current V for current action
            if(tmp[action]>current_max_v):# judge whether it is the maximum V
→value for this state until now
                current_max_v=tmp[action] # If so, update ir
                new_policy[state]=action # If so, update policy

for i in new_policy: # If there exists one state that has no optimal policy
→then break and return False!
    if new_policy[i]==0:
        policy_stable=False

```

```

        break

    # YOUR CODE ENDS HERE
    #####

    return new_policy, policy_stable

```

```

[0]: ## Testing out policy iteration policy for one episode
GAME = "Taxi-v3"
evaluation_utils.render_episode(env, policy_vi)
env = gym.make("Taxi-v3")
V_pi, policy_pi = policy_iteration(env, gamma=0.95, max_iteration=6000,
    ↪theta=1e-5)

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

(North)

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

(West)

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

(West)

```

+-----+
|R: | : :G|

```

```

| : | : : | |
| : | : : |
| | : | : |
|Y| : |B: |

```

```

+-----+

```

(West)

```

+-----+

```

```

|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |

```

```

+-----+

```

(North)

```

+-----+

```

```

|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |

```

```

+-----+

```

(North)

```

+-----+

```

```

|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |

```

```

+-----+

```

(West)

```

+-----+

```

```

|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |

```

```

+-----+

```

(Pickup)

```

+-----+

```

```

|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |

```

```

+-----+

```

(South)

```

+-----+

```

```

|R: | : :G|

```

```

| : | : : |
|█: : : : |
| | : | : |
|Y| : |B: |

```

```
+-----+
```

(South)

```
+-----+
```

```

|R: | : :G|
| : | : : |
| : █: : : |
| | : | : |
|Y| : |B: |

```

```
+-----+
```

(East)

```
+-----+
```

```

|R: | : :G|
| : | : : |
| : █: : : |
| | : | : |
|Y| : |B: |

```

```
+-----+
```

(East)

```
+-----+
```

```

|R: | : :G|
| : |█: : : |
| : : : : |
| | : | : |
|Y| : |B: |

```

```
+-----+
```

(North)

```
+-----+
```

```

|R: |█: :G|
| : | : : : |
| : : : : |
| | : | : |
|Y| : |B: |

```

```
+-----+
```

(North)

```
+-----+
```

```

|R: | : █:G|
| : | : : : |
| : : : : |
| | : | : |
|Y| : |B: |

```

```
+-----+
```

(East)

```
+-----+
```

```

|R: | : :█|

```

```
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
```

(East)

```
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
```

(Dropoff)

Episode reward: 4.000000

```
[0]: # visualize how the agent performs with the policy generated from policy_
      ↪ iteration
      evaluation_utils.render_episode(env, policy_pi)
```

```
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
```

```
+-----+
|R: | : :G| |
| : | : : |
| : | : : |
| | : | : |
|Y| : |B: |
+-----+
```

(South)

```
+-----+
|R: | : :G| |
| : | : : |
| : | : : |
| | : | : |
|Y| : |B: |
+-----+
```

(West)

```
+-----+
|R: | : :G|
| : | : : |
```

```

| : : : : |
|  | : | : |
|Y| : |B: |
+-----+
(South)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(South)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(Pickup)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
|  | : | : |
|Y| : |B: |
+-----+
(North)
+-----+
|R: | : :G| |
| : | : : |
|  | : : : |
| | : | : |
|Y| : |B: |
+-----+
(North)
+-----+
|R: | : :G|
| : | : : |

```

```

| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(North)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(Dropoff)
Episode reward: 11.000000

```

```
[0]: # evaluate the performance of policy iteration over 100 episodes
print(evaluation_utils.avg_performance(env, policy_pi))
```

8.272727272727273

4 Part 3: Q-learning and SARSA

4.1 3.1 Model-free RL: Q-learning

In this part, you will implement Q-learning.

```
[0]: def QLearning(env, num_episodes, gamma, lr, e):
    """Implement the Q-learning algorithm following the epsilon-greedy
    ↪ exploration.
    Update Q at the end of every episode.

    Parameters
    -----
    env: gym.core.Environment
        Environment to compute Q function
    num_episodes: int
        Number of episodes of training.
    gamma: float
        Discount factor.
    learning_rate: float
        Learning rate.
    e: float
        Epsilon value used in the epsilon-greedy method.

    Returns
    -----
    np.array
```



```

An array of shape [env.nS x env.nA] representing state, action values
"""

#####
# YOUR CODE STARTS HERE
Q=np.zeros((env.nS,env.nA)) # initialize every state's Q value to 0
Q_rewards=[] # store cumulative reward for each step

for i in range(num_episodes):
    state=env.reset() # reset the environment
    reward_for_this=0
    while True:
        random_value=np.random.random()
        # epsilon-greedy policy
        if random_value>e:
            action=np.argmax(Q[state]) # choose the maximum value
        else:
            action=np.random.randint(env.nA)
        state_n, reward, done, _ = env.step(action)
        reward_for_this+=reward
        # for q learning, directly choose the action that has the maximum reward
        Q[state][action]=Q[state][action]+lr*(reward+gamma*np.
→max(Q[state_n])-Q[state][action])
        state=state_n # update current state
        if done: # exit if one episode ends, break!!!!
            break
        Q_rewards.append(reward_for_this)
    # YOUR CODE ENDS HERE
#####

#return np.zeros((env.nS, env.nA))
return Q,Q_rewards

```

```

[11]: env = gym.make("Taxi-v3")
Q,Q_rewards = QLearning(env = env.env, num_episodes = 1000, gamma = 1, lr = 0.
→1, e = 0.1)
print('Action values:')
print(Q)

import matplotlib.pyplot as plt
%matplotlib inline
plt.title('Q Learning')
plt.xlabel('Episode')
plt.ylabel('Cumulative Reward of Episode')
plt.plot(range(1,1001),Q_rewards)

```

Action values:


```
[[ 0.          0.          0.          0.          0.          0.          ]]
```

```

[-1.47958384 -3.89193768 -4.20760934 -3.35222425  9.72012492 -4.4645559 ]
[-1.9794029  -0.4017635  -1.84399352 -1.7998      14.2203126  -2.97189524]
...
[-1.1          -1.00877608 -1.1          0.44346311 -1.93         -2.         ]
[-2.88024057 -2.85582085 -2.8924057  -2.26924164 -3.91231215 -6.03047715]
[ 0.37397956 -0.2          -0.2          7.78581602 -1.70920151 -1.         ]]

```

[11]: [<matplotlib.lines.Line2D at 0x7fce1858c160>]



```

[13]: # Uncomment the following to evaluate your result, comment them when you
      ↪ generate the pdf
import evaluation_utils
from utils import action_selection
env = gym.make('Taxi-v3')
policy_estimate = action_selection(Q)
#render(env, policy_estimate)
evaluation_utils.render_episode(env,policy_estimate)

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |

```

+-----+

+-----+

```
|R: | : :G| |
| : | : : |
| : | : : |
| | : | : |
|Y| : |B: |
```

+-----+

(North)

+-----+

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

+-----+

(North)

+-----+

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

+-----+

(North)

+-----+

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

+-----+

(West)

+-----+

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

+-----+

(Pickup)

+-----+

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```

+-----+
  (South)
+-----+
|R: | : :G|
| : | : : |
|█: : : : |
| | : | : |
|Y| : |B: |
+-----+
  (South)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
|█| : | : |
|Y| : |B: |
+-----+
  (South)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
  (South)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
  (Dropoff)
Episode reward: 11.000000

```

4.2 3.2 Model-free RL: SARSA

In this part, you will implement Sarsa.

```

[0]: def SARSA(env, num_episodes, gamma, lr, e):
      """Implement the SARSA algorithm following epsilon-greedy exploration.
      Update Q at the end of every episode.

      Parameters
      -----
      env: gym.core.Environment
          Environment to compute Q function


```

```

num_episodes: int
    Number of episodes of training
gamma: float
    Discount factor.
learning_rate: float
    Learning rate.
e: float
    Epsilon value used in the epsilon-greedy method.

Returns
-----
np.array
    An array of shape [env.nS x env.nA] representing state-action values
    """

#####
# YOUR CODE STARTS HERE
Q=np.zeros((env.nS,env.nA)) # initialize every state's Q value for each
→action
Q_rewards=[]
for i in range(num_episodes):
    state=env.reset() # reset the environment
    # epsilon-greedy policy
    if np.random.random()>e:
        action=np.argmax(Q[state])
    else:
        action=np.random.randint(env.nA)
    reward_for_this=0
    # for SARSA, when choosing next step's action, following the policy of
→current action
    # epsilon-greedy policy
    while True:
        state_n,reward,done,_=env.step(action)
        if np.random.random()>e:
            action_n=np.argmax(Q[state_n])
        else:
            action_n=np.random.randint(env.nA)
        # update Q[state][action]
        □
→Q[state][action]=Q[state][action]+lr*(reward+gamma*Q[state_n][action_n]-Q[state][action])
        reward_for_this+=reward
        state=state_n # update state
        action=action_n # update action
        if done:
            break
    Q_rewards.append(reward_for_this)

```

```

# YOUR CODE ENDS HERE
#####

#return np.ones((env.nS, env.nA))
return Q, Q_rewards

```

```

[0]: import time
def render_episode_Q(env, Q):
    """Renders one episode for Q function on environment.

    Parameters
    -----
    env: gym.core.Environment
        Environment to play Q function on.
    Q: np.array of shape [env.nS x env.nA]
        state-action values.
    """

    episode_reward = 0
    state = env.reset()
    done = False
    while not done:
        env.render()
        time.sleep(0.5)
        action = np.argmax(Q[state])
        state, reward, done, _ = env.step(action)
        episode_reward += reward

    print ("Episode reward: %f" %episode_reward)

```

```

[25]: Q, Q_rewards = SARSA(env = env.env, num_episodes = 10000, gamma = 1, lr = 0.1, e_
    ↪ = 0.1)
print('Action values:')
print(Q)

import matplotlib.pyplot as plt
%matplotlib inline
plt.title('SARSA')
plt.xlabel('Episode')
plt.ylabel('Cumulative Reward of Episode')
plt.plot(range(1,10001),Q_rewards)

```

Action values:

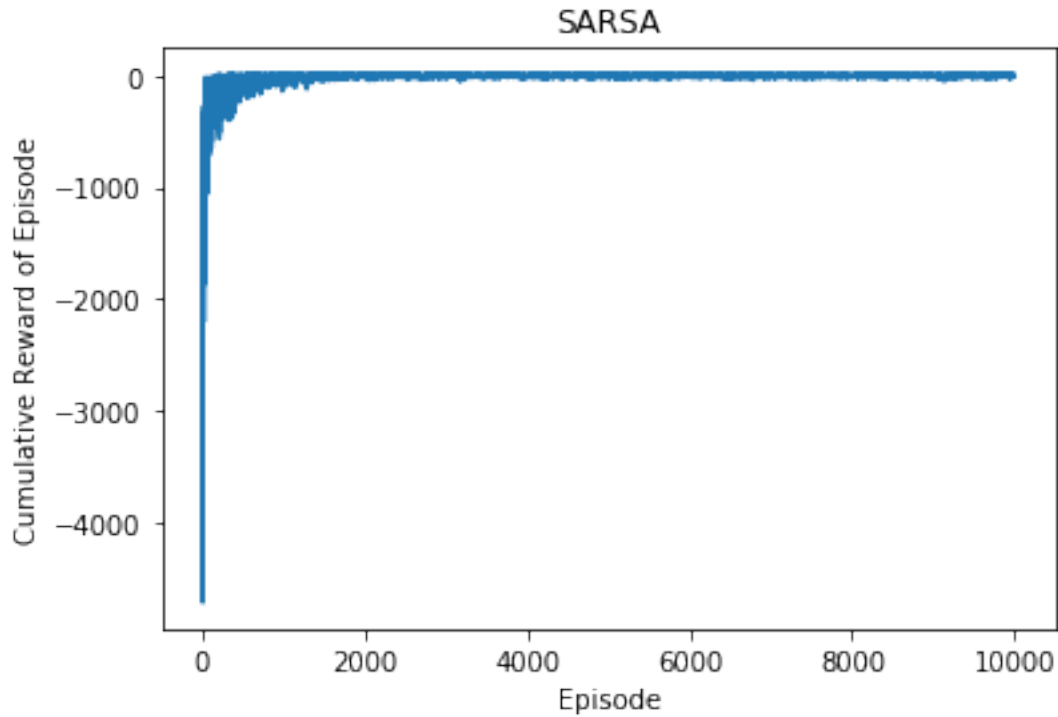
```

[[ 0.          0.          0.          0.          0.          0.          ]
 [ 0.57748287  0.96270149 -1.91369218  2.19509429  6.68944349 -7.37264556]
 [ 4.39720668  7.42154763  5.98620518  6.67607814 12.72281144  0.07580237]
 ...

```

```
[10.02880443 14.29725852  9.05688499  4.85573763  0.64193794 -0.12591817]
[-3.3763171  2.85969628 -3.25709774 -3.66525514 -5.44813421 -6.00780589]
[-0.361      3.96131186  4.44763346 18.73542553 -1.9      -2.57087522]]
```

[25]: [<matplotlib.lines.Line2D at 0x7fce17bcd080>]



```
[26]: # Uncomment the following to evaluate your result, comment them when you
      → generate the pdf
from utils import action_selection
import evaluation_utils
env = gym.make('Taxi-v3')
#policy_estimate = action_selection(Q)
#render(env, policy_estimate)
render_episode_Q(env,Q)
```

```
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

+-----+
```

```

|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(Pickup)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(East)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(South)
+-----+
|R: | : :G| |
| : | : : |
| : | : : |
| | : | : |
|Y| : |B: |
+-----+
(South)
+-----+
|R: | : :G| |
| : | : : |
| : | : : |
| | : | : |
|Y| : |B: |
+-----+
(East)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(North)
+-----+

```



```
|R: | : :G|
| : | : :|
| : : : :|
| | : | :|
|Y| : |B: |
```

```
+-----+
```

(East)

```
+-----+
```

```
|R: | : :G|
| : | : :|
| : : : :|
| | : | :|
|Y| : |B: |
```

```
+-----+
```

(North)

```
+-----+
```

```
|R: | : :G|
| : | : :|
| : : : :|
| | : | :|
|Y| : |B: |
```

```
+-----+
```

(East)

Episode reward: 11.000000